

Car Sales Database Management System

Ajay Vijayakumar, Lokesh Velmurugan Sujatha, Shivaramakrishnan Rajendran
vijayak3@buffalo.edu, lvelmuru@buffalo.edu, rajendr6@buffalo.edu

Abstract—As every industry is moving towards digital era, the auto sales sector is still stuck in the past with old-fashioned methods. There are several difficulties because of this dependency on manual procedures. Dealerships must drop their outdated ledgers and adopt modern technological solutions as the world moves toward a connected, data-driven future.

I. INTRODUCTION

Data Disasters: Imagine you're trying to ask a question or ask for a request online, and the form disappears post you submit it. This is the sad reality of data disasters caused by outdated technology. Missing files and misplaced information are common, resulting in lost sales opportunities and reducing overall efficiency. Important information gets lost, affecting everything from consumer interactions to financial reporting.

Process Pitfalls: that are old fashioned. Compensating for technological deficiencies and dealing with data inaccuracies wastes valuable time and resources. Consider a sales professional attempting to address a customer's question about car availability, and what if we need to get the data from many spreadsheets or handwritten notes. This not only slows down the sales process, but also results in missed opportunities and consumer unhappiness.

Blindfolded Business: Without real-time insights and up to date information, dealerships work without vision. Inventory management becomes a guessing game, with the danger of stock shortages or excesses. Consider a dealership buying more sedans if SUVs are the current fashion and most selling, resulting in unsold stock and reduced business and profits. A lack of visibility into customer trends obstructs strategic decision-making, leaving dealerships unable to respond to market shifts and leverage on the new opportunities.

II. MILESTONE 1

III. DATABASE DESIGN

A. Why do we need a database and how does this solve the problem?

- Structured Data Organization: Imagine having a separate shelf for toys, one for clothes, one for stationaries and so on. This is similar to what a database does. Databases allow structured data storage with separate tables for vehicles, customers, sales, and pricing history. This prevents redundancy and ensures efficient querying.
- Data Relationships and Integrity: In database we can establish relationships between tables which helps us in maintaining the data consistency. For instance, connecting sale to the car and the customer who bought it.

- Granular Permissions and Security: Databases offer fine-grained access control, enhancing security. Depending upon the users authorization and accesses they could just see or edit the data. Real-Time Insights: Tracking sales, inventory, and customer trends in real time empowers managers to make well informed decisions.
- Performance and Efficiency: Databases load faster, consume less storage space, and focus on relational data, laying the foundation for automation.

B. Target User

Consider a streamlined car dealership where information flows smoothly, increasing sales and customer satisfaction. This becomes a reality with the suggested database system, which serves as a central hub for important inventory, sales, and customer history data. Let's look at the important user groups that will unlock its potential:

- **Salespeople:** With the proposed system, employees don't have to rush to find information about the cars or potential buyers. This system provides them with rapid access to car specifications, customer preferences, and previous interactions, helping them to personalize pitches, close deals more quickly, and establish long-term connections.
- **Inventory Managers:** With real-time information on what vehicles are available now, how often a particular car is being sold, and the demand from consumers, they can optimize inventory levels, avoid sales being lost, and make educated purchasing decisions.
- **Marketing Teams:** The proposed database system stores information about the cars which are popular in different regions and customers preferred cars, which helps in targeting the right customers with the accurate offers. They may create targeted campaigns, give personalized promotions, and enhance marketing ROI by leveraging extensive customer data and sales analytics.
- **Management:** From the Management perspective it is easier for them to fix the price for the cars based on the demand it has. They get a complete overview of what the dealership looks like, which helps in making informed decisions about whether to add more employees for making more inventories and to provide promotional offers to the targeted customers.

C. ER Diagram (Before Normalization)

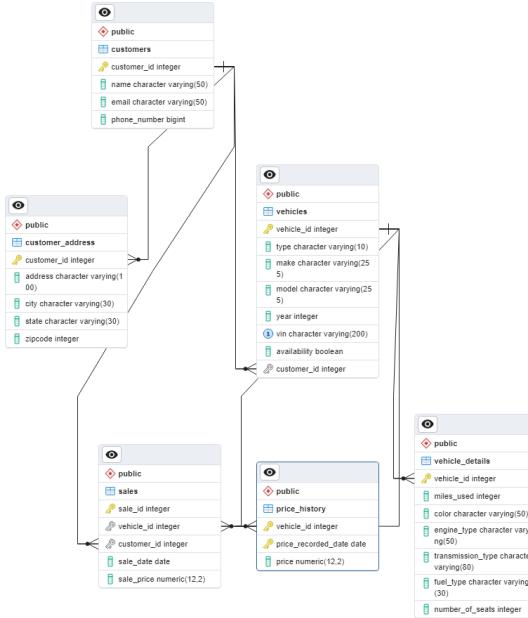


Fig. 1. ER Diagram

D. TABLE

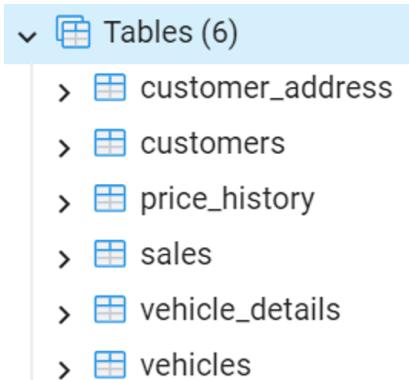


Fig. 2. Tables

| Field | Description |
|--------------|--------------------------------------|
| Customer_id | Unique identifier for each customers |
| Name | Customer's name |
| email | Customer's mail ID |
| phone_number | Customer's phone number |

TABLE I
CUSTOMER TABLE SCHEMA

| Field | Description |
|-------------|---|
| Customer_id | Unique identifier for each customers |
| address | address of the customer |
| city | City where the customer is from |
| state | state where the customer resides |
| zipcode | The postal code of the customer's address |

TABLE II
CUSTOMER_ADDRESS TABLE SCHEMA

| Field | Description |
|--------------|--|
| vehicle_id | Unique identifier for each customers |
| type | vehicle type |
| make | make of the vehicle |
| model | model of the vehicle |
| year | year when vehicle is manufactured |
| vin | Vehicle Identification Number |
| availability | availability of the vehicle |
| customer_id | customer_id of the person owns the vehicle |

TABLE III
VEHICLE TABLE SCHEMA

| Field | Description |
|------------|------------------------------------|
| sale_id | saleid of the transaction |
| vehicle_id | ID of the vehicle purchased |
| sale_date | Date when the transaction happened |
| sale_price | price of the vehicle |

TABLE IV
SALES TABLE SCHEMA

| Field | Description |
|---------------------|-------------------------------------|
| vehicle_id | ID of the vehicle |
| price_recorded_date | Date when the price change happened |
| price | price it changed to |

TABLE V
PRICE_HISTORY TABLE SCHEMA

| | |
|-------------------|--------------------------------------|
| vehicle_id | ID of the vehicle |
| miles_used | miles that the vehicle has travelled |
| color | color of the vehicle |
| engine_type | engine_type of the vehicle |
| transmission_type | transmission_type of the vehicle |
| fuel_type | fuel_type of the vehicle |
| number_of_seats | number_of_seats of the vehicle |

TABLE VI
VEHICLE_DETAILS TABLE SCHEMA

IV. MILESTONE 2

In milestone 2, we focused on improving the database's performance and storage efficiency. We achieved this by refining the initial schema through a process called decomposition. This involved breaking down tables to achieve a specific design standard called Boyce_Codd Normal Form (BCNF).

Specific improvements: The **customer_address** table was split into two separate tables: **customer_address** and **zipcodes**. The **vehicles** table was divided into **vehicles** and

vehicle_models. The **vehicle_details** table was decomposed into **vehicle_details** and **engines**.

V. NORMALIZATION

Database normalization ensures efficient data organization. It resembles organizing belongings for easy retrieval, eliminating duplicates. By dividing data into smaller, logically connected parts, information remains organized, preventing inconsistencies and facilitating updates. This process creates a well-structured filing system for the database. We decompose the table to be in BCNF form.

VI. CONDITIONS FOR BCNF

For a table to meet BCNF, it must pass two tests.

- First, any attribute that determines the value of another must uniquely identify each row. We call these determinants "candidate keys."
- Secondly, if an attribute determines another attribute in a non-trivial manner, it must itself be unique and capable of identifying a row independently. So, when we're looking at each table, we're basically checking if these two rules are satisfied. If they are, then the table is BCNFcompliant. If not, we might need to make some adjustments to ensure everything fits the criteria.

So, when we're looking at each table, we're basically checking if these two rules are satisfied. If they are, then the table is BCNFcompliant. If not, we might need to make some adjustments to ensure everything fits the criteria. To ensure the database adheres to BoyceCodd Normal Form (BCNF), we will begin by identifying non-trivial functional dependencies within each of the six tables. These dependencies will guide us in decomposing tables that violate BCNF principles.

Now Let's look at the current relations and the functional dependencies of those.

A. Table Customers

`customer_id -> name, email, phone_number`
`email -> name`
`phone_number -> name`

The customers table contains information about customers, with a primary key `customer_id`. Each customer has attributes such as `name`, `email`, and `phone_number`, all of which are fully functionally dependent on the primary key. There are no partial dependencies, so the table is in BCNF.

B. Table Customer Address

`customer_id -> address, city, state, zipcode`
`address, city, state -> zipcode (Not in BCNF)`
`zipcode -> state (Not in BCNF)`

The `customer_address` table contains the information such as `customer ID`, `address`, `city`, `state`, and `zipcode`. However, it's

not in Boyce-Codd Normal Form (BCNF) since it does not satisfy the two conditions that we mentioned above. To adhere to BCNF, we decompose the table into two relations. We can figure out that the state just by knowing the zipcode, which means some data is redundant. Hence, we split the table into two.

- In the first table, we just have zipcode and state. Each zipcode corresponds to one state, so there's no redundancy there. This relation is named `zipcodes`.
- In the second table, we keep the rest of the information: zipcode, address, city, and customer ID. Since each customer ID is unique and every combination of address and city is unique (two different customers can't have the same address in the same city), this arrangement avoids redundancy too. This FD confirms to BCNF and there are no more FD that is derived from the parent relation. This relation is named `customer_address`. For this relation, we will have the zipcode as the foreign key.

C. Table Vehicles

`vehicle_id -> type, make, model, year, vin,`
`availability, customer_id`
`vin -> type, make, model, year`
`model -> make (Not in BCNF)`

The vehicles table, containing attributes like vehicle ID, type, make, model, year, VIN, availability, and customer ID, exhibits functional dependencies which is not in BoyceCodd Normal Form (BCNF). To adhere to BCNF, decomposition is necessary. One relation is formed with attributes

`{vehicle_id, type, model, year, vin,`
`availability, customer_id}`

where vehicle ID uniquely determines these attributes and the model is a foreign key. Another relation includes model, make is created to resolve the dependency between model and make. This decomposition ensures each relation adheres to BCNF, eliminating redundancy and maintaining data integrity within the vehicles dataset.

D. Table Vehicle Details

`vehicle_id -> miles_used, color ,`
`engine_type, transmission_type,`
`fuel_type, number_of_seats`
`engine_type -> fuel_type`

E. Table Sales

`sale_id -> vehicle_id, customer_id,`
`sale_date, sale_price`
`vehicle_id -> customer_id, sale_date,`
`sale_price`

The table `vehicle_details` serves as a repository for essential information pertaining to vehicles, encompassing attributes such as vehicle ID, miles used, color, engine type, transmission type, fuel type, and number of seats. As we delved into the dataset, we discerned a significant concern regarding the relationship between engine type and fuel type. Considering this observation, we recognized the necessity of reorganizing the dataset to address this issue. Our aim was to enhance the structure and coherence of the data, ensuring that it aligns with established standards for data organization and normalization. To achieve this objective, we do decomposition, aimed at restructuring the "vehicle_details" table into more streamlined and logically coherent components. We split the table into two relations: The first one has the attributes (`engine_type`, `fuel_type`) and the second one with all the other attributes. These relations are in BCNF, and no further decomposition is required. The first relation is named `engines`, and second relation is named `vehicle_details` in which the `engine_type` is a foreign key.

F. Table Price_History

`vehicle_id, price_recorded_date -> price`

The "price_history" table keeps track of how prices for vehicles change over time. Each entry in the table includes the vehicle ID, the date when the price was recorded, and the price itself. This setup follows a rule called Boyce-Codd Normal Form (BCNF), which means it's organized in a way that makes sense. Each combination of vehicle ID and recording date uniquely determines the price at that time, without any unnecessary connections or complications. So, the "price_history" table is already structured in a way that fits the BCNF rules.

VII. DATABASE DESIGN - ER DIAGRAM (AFTER NORMALIZATION)



Fig. 3. ER Diagram after normalization

VIII. TABLES (AFTER NORMALIZATION)

| Tables (9) |
|--------------------|
| > customer_address |
| > customers |
| > engines |
| > price_history |
| > sales |
| > vehicle_details |
| > vehicle_models |
| > vehicles |
| > zipcodes |

Fig. 4. TABLES After Normalization

| Field | Description |
|--------------|--------------------------------------|
| Customer_id | Unique identifier for each customers |
| Name | Customer's name |
| email | Customer's mail ID |
| phone_number | Customer's phone number |

TABLE VII
CUSTOMER TABLE SCHEMA

| Field | Description |
|-------------|---|
| Customer_id | Unique identifier for each customers |
| address | address of the customer |
| city | City where the customer is from |
| zipcode | The postal code of the customer's address |

TABLE VIII
CUSTOMER_ADDRESS TABLE SCHEMA

| Field | Description |
|---------|---|
| zipcode | The postal code of the customer's address |
| state | state where the customer resides |

TABLE IX
ZIPCODES TABLE SCHEMA

| Field | Description |
|--------------|--|
| vehicle_id | Unique identifier for each customers |
| type | vehicle type |
| model | model of the vehicle |
| year | year when vehicle is manufactured |
| vin | Vehicle Identification Number |
| availability | availability of the vehicle |
| customer_id | customer_id of the person owns the vehicle |

TABLE X
VEHICLE TABLE SCHEMA

| Field | Description |
|-------|----------------------|
| model | model of the vehicle |
| make | make of the vehicle |

TABLE XI
MODEL TABLE SCHEMA

| Field | Description |
|------------|------------------------------------|
| saleid | saleid of the transaction |
| vehicle_id | ID of the vehicle purchased |
| sale_date | Date when the transaction happened |
| sale_price | price of the vehicle |

TABLE XII
SALES TABLE SCHEMA

| Field | Description |
|---------------------|-------------------------------------|
| vehicle_id | ID of the vehicle |
| price_recorded_date | Date when the price change happened |
| price | price it changed to |

TABLE XIII
PRICE_HISTORY TABLE SCHEMA

| | |
|-------------------|--------------------------------------|
| vehicle_id | ID of the vehicle |
| miles_used | miles that the vehicle has travelled |
| color | color of the vehicle |
| engine_type | engine_type of the vehicle |
| transmission_type | transmission_type of the vehicle |
| number_of_seats | number_of_seats of the vehicle |

TABLE XIV
VEHICLE_DETAILS TABLE SCHEMA

| Field | Description |
|-------------|----------------------------|
| engine_type | engine_type of the vehicle |
| fuel_type | fuel_type of the vehicle |

TABLE XV
PRICE_HISTORY TABLE SCHEMA

IX. DATABASE LOOKUP

A. Customer Table:

| Query Query History | | | |
|--|--------------|------------------|----------------------------|
| 1 <code>SELECT * FROM customers;</code> | | | |
| Data Output Messages Notifications | | | |
| customer_id | [PK] integer | name | character varying (50) |
| 1 | 866 | Barbara Buchanan | bryan08@example.org |
| 2 | 448 | Andrea Roach | xmorgan@example.org |
| 3 | 822 | James Baxter | rebeccadavis@example.com |
| 4 | 2099 | Michael Johnson | everettcharles@example.com |
| 5 | 706 | Cynthia Mitchell | mmiller@example.net |
| 6 | 429 | Melissa Walker | cobbmelissa@example.com |

B. Customer Address Table:

| Query Query History | | | |
|--|--------------|------------------|----------------------------|
| 1 <code>SELECT * FROM customers;</code> | | | |
| Data Output Messages Notifications | | | |
| customer_id | [PK] integer | name | character varying (50) |
| 1 | 866 | Barbara Buchanan | bryan08@example.org |
| 2 | 448 | Andrea Roach | xmorgan@example.org |
| 3 | 822 | James Baxter | rebeccadavis@example.com |
| 4 | 2099 | Michael Johnson | everettcharles@example.com |
| 5 | 706 | Cynthia Mitchell | mmiller@example.net |
| 6 | 429 | Melissa Walker | cobbmelissa@example.com |

C. Engines Table:

| 1 <code>SELECT * FROM engines;</code> | | | |
|--|-----------------------------|------------------|------------------------|
| Data Output Messages Notifications | | | |
| engine_type | [PK] character varying (50) | fuel_type | character varying (30) |
| 1 | Gas (V6) | Regular unleaded | |
| 2 | Gasoline (V6) | Regular unleaded | |
| 3 | Electric (Battery-powered) | Electricity | |
| 4 | Gasoline (Petrol) | Regular unleaded | |
| 5 | Gas/Electric I-4 | Regular unleaded | |

D. Price History Table:

| Query Query History | | | |
|---|--------------|---------------------|-----------|
| 1 <code>SELECT * FROM price_history;</code> | | | |
| Data Output Messages Notifications | | | |
| vehicle_id | [PK] integer | price_recorded_date | [PK] date |
| 1 | 214363 | 2021-02-21 | 45587.00 |
| 2 | 214363 | 2015-09-26 | 44098.00 |
| 3 | 214363 | 2018-09-25 | 43442.00 |
| 4 | 214363 | 2021-11-10 | 46004.00 |
| 5 | 214363 | 2018-01-29 | 43023.00 |
| 6 | 214363 | 2021-03-29 | 42003.00 |
| 7 | 813950 | 2003-08-11 | 43226.00 |
| 8 | 813950 | 2013-04-20 | 42503.00 |
| 9 | 813950 | 2001-09-04 | 42888.00 |
| 10 | 813950 | 2001-04-23 | 43939.00 |

E. Sales Table:

Query Query History

```
1 SELECT * FROM sales;
```

Data Output Messages Notifications

| | sale_id | vehicle_id | customer_id | sale_date | sale_price |
|----|--------------|------------|-------------|------------|----------------|
| | [PK] integer | integer | integer | date | numeric (12,2) |
| 1 | 956807 | 214363 | 866 | 2023-08-06 | 23414.00 |
| 2 | 289381 | 813950 | 448 | 2013-08-22 | 30958.00 |
| 3 | 784217 | 494481 | 822 | 2019-09-25 | 38387.00 |
| 4 | 633004 | 388014 | 2099 | 2014-11-01 | 24826.00 |
| 5 | 932420 | 287300 | 706 | 2015-09-29 | 25223.00 |
| 6 | 377444 | 852621 | 429 | 2009-09-09 | 41590.00 |
| 7 | 317663 | 686118 | 2297 | 2023-08-28 | 37528.00 |
| 8 | 76327 | 990403 | 3638 | 2023-06-25 | 25861.00 |
| 9 | 809545 | 801366 | 2350 | 2023-05-14 | 56349.00 |
| 10 | 907456 | 517071 | 1156 | 2023-06-27 | 34589.00 |

F. Vehicle Details Table:

Query Query History

```
1 SELECT * FROM vehicle_details;
```

Data Output Messages Notifications

| | vehicle_id | miles_left | color | engine_type | transmission_type | number_of_seats |
|----|--------------|------------|------------------------|----------------------------|---|-----------------|
| | [PK] integer | [null] | character varying (50) | character varying (50) | character varying (80) | integer |
| 1 | 214363 | [null] | Silver | Hybrid (2.5L FHEV) | PowerSplit: Electronic continuously variable transmission (CVT) | 5 |
| 2 | 813950 | [null] | Red | Turbocharged Gas I4 | 8-speed automatic | 5 |
| 3 | 494481 | [null] | White | Gas (V6) | 8-speed shiftable automatic | 7 |
| 4 | 388014 | 27567 | White | Gasoline (I-4) | CVT | 5 |
| 5 | 287300 | [null] | Red | Gasoline (Petrol) | Continuously variable-speed automatic (CVT) | 5 |
| 6 | 852621 | [null] | White | Gas (V6) | 9-speed automatic | 5 |
| 7 | 686118 | \$340 | Grey | Gas (V6) | 8-speed shiftable automatic | 7 |
| 8 | 990403 | [null] | White | Gasoline (V6) | Xtronic CVT | 5 |
| 9 | 801366 | [null] | Red | Electric (Battery-powered) | Automatic | 5 |
| 10 | 517071 | [null] | Red | Gas/Electric I-4 | CVT | 5 |

G. Vehicle Models Table:

Query Query History

```
1 SELECT * FROM vehicle_models;
```

Data Output Messages Notifications

| | model | make |
|----|------------------------------|-------------------------|
| | [PK] character varying (255) | character varying (255) |
| 1 | RAV4 Hybrid | Toyota |
| 2 | Bronco Sport | Ford |
| 3 | Edge | Ford |
| 4 | Equinox | Chevrolet |
| 5 | Mustang Mach-E | Ford |
| 6 | Blazer EV | Chevrolet |
| 7 | Highlander | Toyota |
| 8 | Tundra | Toyota |
| 9 | F-150 | Ford |
| 10 | Rouge | Nissan |

H. Vehicles Table:

Query Query History

```
1 SELECT * FROM vehicles;
```

Data Output Messages Notifications

| | vehicle_id | type | model | year | vin | availability | customer_id |
|----|--------------|------------------------|-------------------------|---------|---|--------------|-------------|
| | [PK] integer | character varying (10) | character varying (255) | integer | character varying (200) | boolean | integer |
| 1 | 214363 | new | Maverick | 2021 | aebdb39b-378f-4254-8199-934a1afab... | false | 866 |
| 2 | 813950 | new | Colorado | 2020 | 0dab04c4-289f-4ed8-bc1d-92595617... | false | 448 |
| 3 | 188383 | new | F-150 | 2007 | 5c01cc09-59c9-4c46-b033-9556e9391b... | true | [null] |
| 4 | 494481 | new | Highlander | 2006 | 8c42eff7-7ad4-a8a1-536d7a1a1844 | false | 822 |
| 5 | 388014 | used | HR-V | 2004 | 2d3b389f-d131-4bf5-b5fa-9cc0cc5cb523 | false | 2099 |
| 6 | 287300 | new | Corolla | 2005 | 15256feef2f1-40cd-9a99-425b5af997 | false | 706 |
| 7 | 852621 | new | Ridgeline | 2001 | 6a9cf3a7-ced8-4ea4-9d1a-9f94a252cff | false | 429 |
| 8 | 686118 | used | Highlander | 2017 | c172e071-9a64-4814-948e-e439a2ee202c... | false | 2297 |
| 9 | 990403 | new | Maxima | 2018 | f119a2b-d6f5-43c6-880a-47bce7257e | false | 3638 |
| 10 | 801366 | new | Blazer EV | 2021 | ed1ecea4-c032-45b6-af5b-035cdd8d08... | false | 2350 |

I. Zipcodes Table:

Query Query History

```
1 SELECT * FROM zipcodes;
```

Data Output Messages Notifications

| | zipcode | state |
|---|----------------|------------------------|
| | [PK] character | character varying (30) |
| 1 | 00601 | Puerto Rico |
| 2 | 00624 | Puerto Rico |
| 3 | 00662 | Puerto Rico |
| 4 | 00676 | Puerto Rico |
| 5 | 00685 | Puerto Rico |
| 6 | 00690 | Puerto Rico |

X. QUERY EXECUTION

A. Find the distribution of the vehicle sales among different car models.

Query Query History

```
--#1 The distribution of vehicle sales among different car models
1
2
3
4
5
6
7
8
```

Data Output Messages Notifications

| | model | vehicles_sold |
|----|-------------------------|---------------|
| | character varying (255) | bigint |
| 1 | Corolla | 328 |
| 2 | Maverick | 324 |
| 3 | Passport | 303 |
| 4 | Highlander | 284 |
| 5 | Trax | 275 |
| 6 | RAV4 Hybrid | 256 |
| 7 | F-150 | 239 |
| 8 | Silverado | 239 |
| 9 | CR-V Hybrid | 227 |
| 10 | Tundra | 222 |

B. Find all the customers who purchased more than two vehicles.

```
Query  Query History
1 --#2 Customers who purchased more than two vehicles
2
3 WITH vehicles_per_cust AS
4   (SELECT customer_id,
5    count(vehicle_id) AS vehicles_count
6   FROM sales
7   GROUP BY customer_id
8   HAVING count(vehicle_id) > 2)
9 SELECT name,
10       email,
11       vehicles_count
12  FROM vehicles_per_cust
13 JOIN customers cust ON cust.customer_id = vehicles_per_cust.customer_id;
```

Data Output Messages Notifications

| | name | email | vehicles_count |
|----|-------------------|--------------------------|----------------|
| 1 | Theresa Woods | wgordon@example.com | 3 |
| 2 | Alex Phillips | vjenkins@example.org | 3 |
| 3 | Steven Munoz | duncarjudith@example.org | 3 |
| 4 | Shannon Frost | dalvakristen@example.net | 3 |
| 5 | Elizabeth Johnson | kellyonitz@example.net | 3 |
| 6 | Nicole Cohen | ygonzalez@example.com | 3 |
| 7 | Jennifer Peterson | ivvaughan@example.com | 3 |
| 8 | Gregory Moore | katelyn35@example.com | 3 |
| 9 | Miguel Henry | mgraham@example.org | 3 |
| 10 | Tiffany Walker | cmason@example.net | 3 |

C. Find the average number of days between consecutive vehicle purchases for customers who have bought more than one vehicle.

```
Query  Query History
1 --#3 Average number of days between consecutive vehicle purchases for
2 -- customers who have bought more than one vehicle?
3
4 WITH vehicles_per_cust AS
5   (SELECT customer_id,
6    count(vehicle_id) AS vehicles_count
7   FROM sales
8   GROUP BY customer_id
9   HAVING count(vehicle_id) > 1),
10  days_bwn_purchase AS
11  (SELECT vehicles_per_cust.customer_id,
12    sale_date,
13    lag(sale_date) OVER (PARTITION BY vehicles_per_cust.customer_id
14                           ORDER BY sale_date) AS previous_purchase_date
15  FROM vehicles_per_cust
16  JOIN sales ON vehicles_per_cust.customer_id = sales.customer_id)
17  SELECT round(avg(sale_date - previous_purchase_date)) AS avg_days_between_purchase
18  FROM days_bwn_purchase
19 WHERE previous_purchase_date IS NOT NULL;
```

Data Output Messages Notifications

| | avg_days_between_purchase |
|---|---------------------------|
| 1 | 1860 |

D. Find the top 5 cars with the highest sales volume.

```
Query  Query History
1 --#4 Top 5 cars with the highest sales volume
2 WITH sales_volume AS
3   (SELECT vehicle_models.model,
4    vehicle_models.make,
5    COUNT(*) AS sales_volume
6   FROM vehicles
7   JOIN vehicle_models ON vehicles.model = vehicle_models.model
8   JOIN sales ON vehicles.vehicle_id = sales.vehicle_id
9   GROUP BY vehicle_models.model,
10    vehicle_models.make
11   ORDER BY sales_volume DESC)
12  SELECT model,
13        make
14  FROM sales_volume
15  LIMIT 5;
```

Data Output Messages Notifications

| model | make |
|--------------|-----------|
| 1 Corolla | Toyota |
| 2 Maverick | Ford |
| 3 Passport | Honda |
| 4 Highlander | Toyota |
| 5 Trax | Chevrolet |

E. Find the models with highest sales in each state.

```
Query  Query History
1 --#5 Model with the highest sales in each state
2 WITH sales_volume_by_state AS
3   (SELECT model,
4    state,
5    COUNT(model) AS model_count,
6    RANK() OVER (PARTITION BY state
7                  ORDER BY COUNT(model) DESC) AS sales_rank
8   FROM sales
9   JOIN customer_address ON sales.customer_id = customer_address.customer_id
10  JOIN zipcodes ON customer_address.zipcode = zipcodes.zipcode
11  JOIN customers ON customer_address.customer_id = customers.customer_id
12  JOIN vehicles ON customers.customer_id = vehicles.customer_id
13 WHERE availability = FALSE
14 GROUP BY model,
15       state)
16 SELECT state,
17       model
18  FROM sales_volume_by_state
19 WHERE sales_rank = 1;
```

Data Output Messages Notifications

| state | model |
|---------------|------------|
| 1 Alabama | Trax |
| 2 Alaska | Highlander |
| 3 Alaska | Trax |
| 4 Arizona | Maxima |
| 5 Arkansas | Equinox |
| 6 California | Silverado |
| 7 Colorado | Maxima |
| 8 Connecticut | Maverick |
| 9 Delaware | Sienna |
| 10 Delaware | Equinox |

INSERTION

F. Inserting Customer who hasn't made any vehicle purchases.

```
Query  Query History
1 --#6 Add a customer who hasn't made any vehicle purchases
2 INSERT INTO customers(customer_id, name, email, phone_number)
3 VALUES (4000,'Elijah Evans','elijahfevans@example.net',9177523920),
4        (4001,'James Case','jamesacase@example.com',5024097921);
```

Data Output Messages Notifications

INSERT 0 2

Query returned successfully in 56 msec.

UPDATION

G. Updating the Customers address.

```
Query  Query History
1 --#7 Update customer address
2 UPDATE customer_address
3 SET address = '155 Pitts Lodge'
4 WHERE customer_id = 2582;
```

Data Output Messages Notifications

UPDATE 1

Query returned successfully in 37 msec.

DELETION

H. Deleting the incorrect price from price history.

Query Query History

```
1 --#8 Delete incorrect price from price history
2 DELETE
3 FROM price_history
4 WHERE vehicle_id = 7511
5 AND price_recorded_date='2009-04-25';
```

Data Output Messages Notifications

DELETE 1

Query returned successfully in 52 msec.

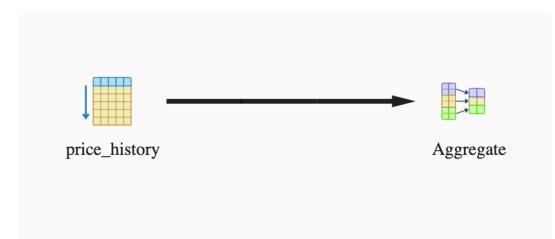
XI. INDEXING

While creating the database and running the SQL queries we found that a few of the queries were running for a long time. Upon investigating this we found that the queries were doing a full table scan before joining them. This can be rectified by creating and using indexes on tables.

Indexes enable faster access to data in a large table. Indexes enhance performance, especially in large data sets, when columns involved in querying large data sets are indexed. For example, if the primary key or filtering/ordering columns are often used to retrieve data from the database, the DBMS can quickly find the rows relevant to the query, without having to read every row of the table. Queries with filtering, sorting, and joining data based on the indexed columns can have significant performance improvement. The DBMS will be able to find the relevant rows quickly using the index rather than full-table scans for querying, which would increase the time for query execution.

We have created five indexes namely: An index named **vehicles_availability_idx** is used to enhance the retrieval of the vehicle availability status thereby, when needed, it is possible to check whether a vehicle is available or not in a matter of seconds. To support faster querying of vehicle models, because of the frequent insertion into WHERE clauses, **vehicles_model_idx** is designed. An index named **vehicles_type_idx** is designed to retrieve a user preference for either new or pre-owned vehicles quickly. An index **price_history_price_idx** is created for querying the price. An index named **price_history_price_idx** is established to optimize querying of price ranges within WHERE conditions, recognizing its frequent usage. e range in where condition, this is one of the most used conditions.

A. Without-Indexing (Price between range)



The diagram illustrates a full table scan process. On the left, a grid icon with a downward arrow is labeled "price_history". A thick black arrow points from this to the right, where another grid icon with a plus sign is labeled "Aggregate".

Query Query History

```
1 --#9 Price between range
2 SELECT min(price) AS lowest,
3        max(price) AS highest
4 FROM price_history
5 WHERE price BETWEEN 41650 AND 41750;
```

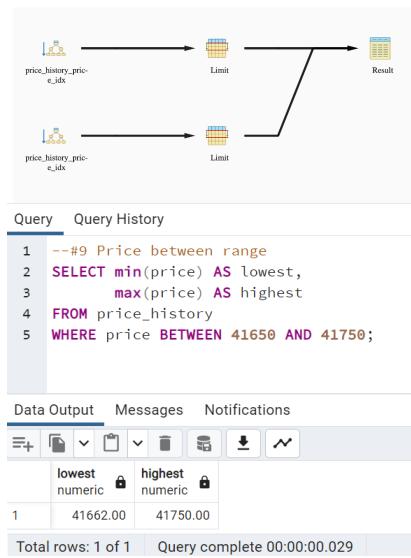
Data Output Messages Explain Notifications

| lowest | highest |
|---------|----------|
| numeric | numeric |
| 1 | 41662.00 |
| | 41750.00 |

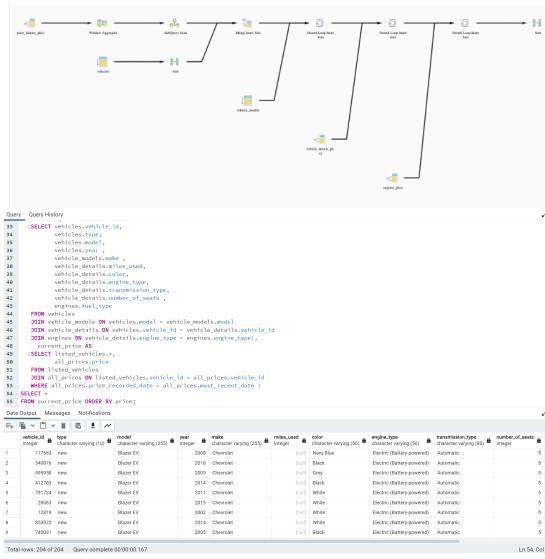
Total rows: 1 of 1 Query complete 00:00:00.059

We have used B+ indexing for finding the price between ranges. B+ indexing facilitates efficient retrieval of data within a given price range. The prices are sorted within a balanced tree so that the prices of a particular category are organized within a balanced tree. Each node within the tree points to child nodes or to data records, allowing quick navigation through subsets of data that meet a condition given. For example, a query requesting a range of prices is efficiently done through the B+ tree since it identifies the leaf nodes that contain price records in the specified range. It does this through its sorted order and balanced height properties in identifying the proper data that must be accessed to ensure that the query execution time is minimal, yet the performance is enhanced for the price range conditions.

B. With-Indexing (Price between range):



C. Without-Indexing:

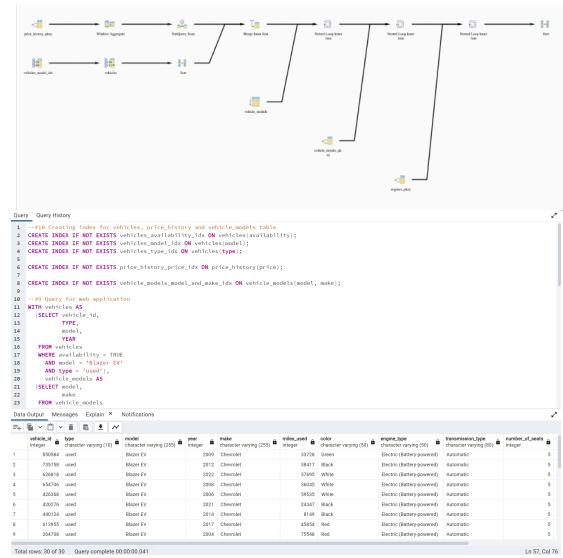


This index is very important in optimizing model-related queries. This index will allow the database system to directly move to the specific model being queried without having to scan the entire vehicles table. Therefore, the index makes it possible for the system to directly fetch the rows that are associated with the model being queried without performing a costly full-table scan.

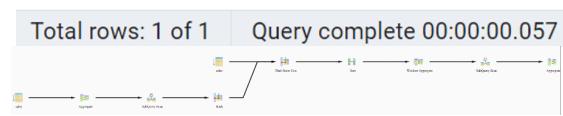
Observations indicate that indexing significantly improves performance, the performance increased around 4 times as compared to the one without indexing as evidenced by notably reduced retrieval times compared to operations conducted without indexing.

Similarly for the query where we found the average number of days between consecutive vehicle purchases for customers who have bought more than one vehicle, we have indexed the sales table to improve the query performance.

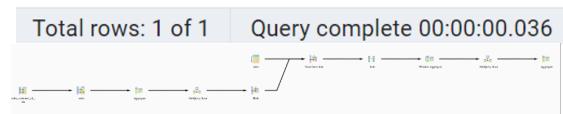
D. With Indexing:



E. Without-Indexing:



F. With Indexing:



XII. WEBPAGE

The website is hosted in streamlit community servers, and the database is hosted in AWS RDS instances. The website is accessible through the following link:
<https://cse-560-dmql-proj-rangers.streamlit.app/>

This page contains the page for buying cars, price history of the car model, and the cars that are sold. In the page buy car it contains the inputs that is got from the user such as Make, Model, Type, Sort By (Price).

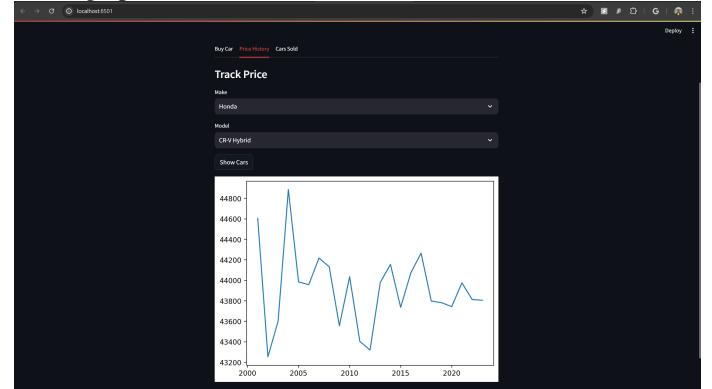
| | vehicle_id | type | model | year | make | miles_used | color | engine_type | transmissio |
|---|------------|------|-------------|-------|-------|------------|-----------|------------------|-------------|
| 0 | 729,961 | new | CR-V Hybrid | 2,001 | Honda | None | Black | Gas/Electric I-4 | CVT |
| 1 | 771,901 | new | CR-V Hybrid | 2,006 | Honda | None | Grey | Gas/Electric I-4 | CVT |
| 2 | 202,003 | new | CR-V Hybrid | 2,020 | Honda | None | Black | Gas/Electric I-4 | CVT |
| 3 | 353,213 | new | CR-V Hybrid | 2,012 | Honda | None | White | Gas/Electric I-4 | CVT |
| 4 | 190,974 | new | CR-V Hybrid | 2,001 | Honda | None | White | Gas/Electric I-4 | CVT |
| 5 | 698,747 | new | CR-V Hybrid | 2,011 | Honda | None | Black | Gas/Electric I-4 | CVT |
| 6 | 351,377 | new | CR-V Hybrid | 2,006 | Honda | None | Navy Blue | Gas/Electric I-4 | CVT |
| 7 | 477,952 | new | CR-V Hybrid | 2,017 | Honda | None | Red | Gas/Electric I-4 | CVT |
| 8 | 825,383 | new | CR-V Hybrid | 2,010 | Honda | None | Grey | Gas/Electric I-4 | CVT |
| 9 | 349,066 | new | CR-V Hybrid | 2,008 | Honda | None | Red | Gas/Electric I-4 | CVT |

If a customer is looking to purchase a new Honda CR-V Hybrid and want to ensure you're getting a competitive price, the Autocars website offers a streamlined process to assist you. By selecting Honda as the make and CR-V Hybrid as the model, and specifying the type as new, you can sort the available vehicles by price. This feature allows you to view a list of new Honda CR-V Hybrids in ascending order of price, making it easier to find the best deal within your budget. Each listing provides comprehensive details such as the vehicle ID, model year, color, engine type, and transmission, ensuring you have all the information needed to make an informed decision.

| | vehicle_id | type | model | year | make | miles_used | color | engine_type | transmissio |
|---|------------|------|-------------|-------|-------|------------|-----------|------------------|-------------|
| 0 | 729,961 | new | CR-V Hybrid | 2,001 | Honda | None | Black | Gas/Electric I-4 | CVT |
| 1 | 771,901 | new | CR-V Hybrid | 2,006 | Honda | None | Grey | Gas/Electric I-4 | CVT |
| 2 | 202,003 | new | CR-V Hybrid | 2,020 | Honda | None | Black | Gas/Electric I-4 | CVT |
| 3 | 353,213 | new | CR-V Hybrid | 2,012 | Honda | None | White | Gas/Electric I-4 | CVT |
| 4 | 190,974 | new | CR-V Hybrid | 2,001 | Honda | None | White | Gas/Electric I-4 | CVT |
| 5 | 698,747 | new | CR-V Hybrid | 2,011 | Honda | None | Black | Gas/Electric I-4 | CVT |
| 6 | 351,377 | new | CR-V Hybrid | 2,006 | Honda | None | Navy Blue | Gas/Electric I-4 | CVT |
| 7 | 477,952 | new | CR-V Hybrid | 2,017 | Honda | None | Red | Gas/Electric I-4 | CVT |
| 8 | 825,383 | new | CR-V Hybrid | 2,010 | Honda | None | Grey | Gas/Electric I-4 | CVT |
| 9 | 349,066 | new | CR-V Hybrid | 2,008 | Honda | None | Red | Gas/Electric I-4 | CVT |

In the second page of the website, it shows the graph of the price history for a make and model of the car selected by the user. For instance in the above graph it shows price history of the Nissan Altima spanning several years. When users select "Nissan" as the make and "Altima" as the model within the "Price History" tab, the chart indicates the way in which the price fluctuates from year to year. Notice the distinct cyclical trend: the price increases and decreases on an alternating basis. This is a visual insight that allows users to understand market value fluctuations of the Nissan Altima.

The price history page gives us the information of the previous price of the make, model and the outputs is given as a line graph.



The third page consist of the map that shows the cars that are sold in a region. By observing the concentration of blue dots, users can infer which car models are popular in different parts of the region. Areas with a high density of sales likely correspond to popular car choices among local buyers. Users can use this map to identify trends and understand which car models are in demand. If certain areas consistently show higher sales density for specific car brands or models, those vehicles are likely popular in that region.

